

98point6 Data Engineer Homework

Welcome to 98point6's at-home interview question! The purpose of this exercise is for you to demonstrate your data engineering competency by designing and implementing a basic data warehouse with a database engine (like Amazon Redshift, PostgreSQL, MySQL, or similar). The data warehouse will be used to understand player and game information for an internal game, 98point6 Drop Token (described in next section).

Your task is to design a data warehouse for stakeholders to answer questions about all the 98point6 Drop Token (9dt) players and games, and generally explore the data. You have access to a CSV file with game data and a player profile API that returns detailed player information based on the player id.

Rules of the 98point6 Drop Token Game

9dt takes place on a 4x4 grid. A token is dropped along a column and said token goes to the lowest unoccupied row of the board. A player wins when they have 4 tokens next to each other either along a row, in a column, or on a diagonal. If the board is filled, and nobody has won then the game is a draw. Each player takes a turn, starting with player 1, until the game reaches either win or draw. If a player tries to put a token in a column that is already full, that results in an error state, and the player must play again until they play a valid move.

Example Game

9dt game data

The game_data.csv file is available at:

https://s3-us-west-2.amazonaws.com/98point6-homework-assets/game_data.csv

It contains four columns with each row indicating a single move by a single player in a game. The columns are:

1. game_id - Opaque string indicating with unique identifier for a specific game that was played between two players.
2. player_id - Opaque string indicating which player made a move.
3. move_number - Integer indicating the move number for this game. The first move of a game will have 1 as the move_number.
4. column - Integer column that the player dropped their token into.
5. result - Contains "win" when the move results in the player winning the game, "draw" if the move ends the game in draw, and is empty otherwise.

1

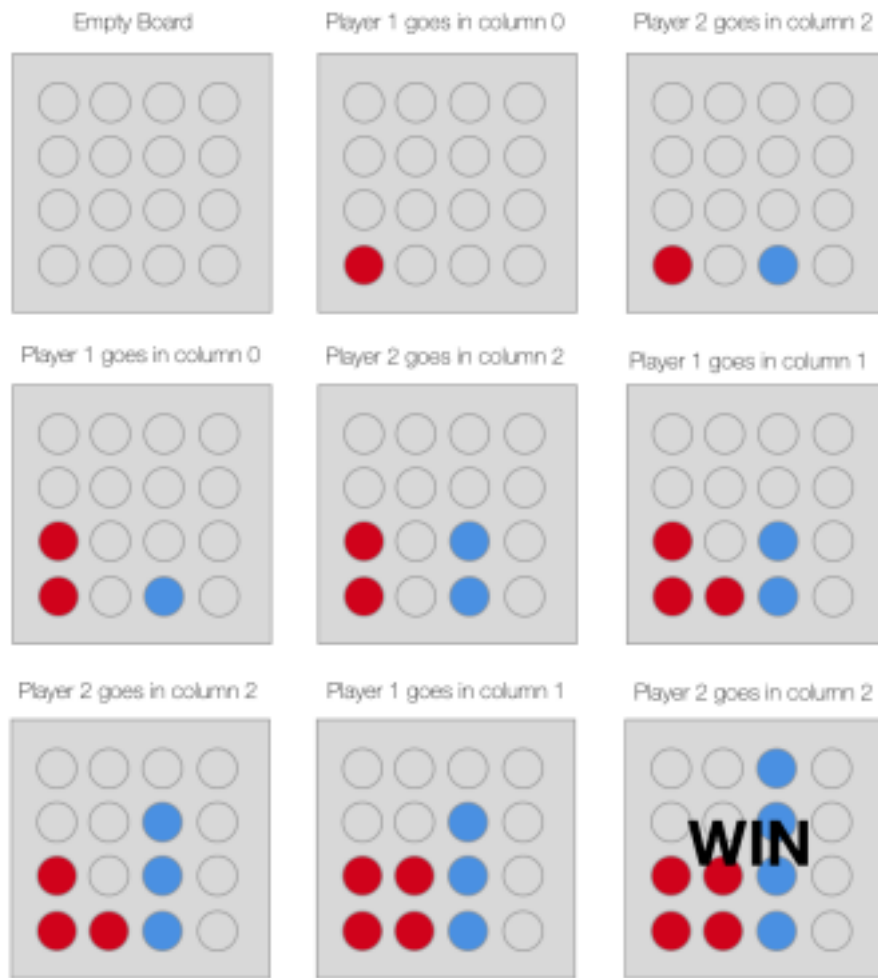


Figure 1: samplegame

2

9dt player profile API

There is a 9dt player profile API at the following location:

<https://x37sv76kth.execute-api.us-west-1.amazonaws.com/prod/users>. You can make HTTP GET requests to the API to retrieve JSON formatted data about players, 10 at a time:

```
curl -X GET -H 'Accept: application/json' \
      https://x37sv76kth.execute-api.us-west-1.amazonaws.com/prod/users?page=0
```

Increase the page query string parameter to get more players. An empty array is returned when there are no more players (at time of writing, there were ~5000 players).

An example response is:

```
[
  {
    "id": 3300,
    "data": {
      "gender": "male",
      "name": {
        "title": "monsieur",
        "first": "adrien",
        "last": "lambert"
      },
      "location": {
        "street": "2520 rue barrier",
        "city": "morges 2",
        "state": "nidwalden",
        "postcode": 9752
      },
      "email": "adrien.lambert@example.com",
      "login": {
        "username": "yellowgorilla717"
      },
      "dob": "1946-01-26 18:44:10",
      "registered": "2007-05-17 07:51:54",
      "phone": "(127)-093-1911",
      "cell": "(867)-527-4578",
      "id": {
        "name": "AVS",
        "value": "756.OUGI.JBUO.16"
      },
      "picture": {
        "large": "https://randomuser.me/api/portraits/men/81.jpg", "medium":
        "https://randomuser.me/api/portraits/med/men/81.jpg", "thumbnail":
        "https://randomuser.me/api/portraits/thumb/men/81.jpg"
      },
      "nat": "CH"
    }
  }
]
```

```

    }
  },
  ...
]

```

The id field at the top-level of each object in the response is the player id (same as player_id in CSV). The data object has details about the player like name, nationality (nat), email, and more (player data from <https://randomuser.me/>).

Analysis

Your data warehouse should enable its consumers to easily run the following analyses:

1. Out of all the games, what is the percentile rank of each column used as the first move in a game? That is, when the first player is choosing a column for their first move, which column most frequently leads to that player winning the game?
2. How many games has each nationality participated in?
3. Marketing wants to send emails to players that have only played a single game. The email will be customized based on whether or not the player won, lost, or drew the game. Which players should receive an email, and with what customization?

There should be tables or views to support each of the above analyses. Please include the table or view names that relate to each of the above analyses in your submission.

Submitting your implementation

We expect one deliverable in your submission with the scripts or executables that setup the data warehouse, load the data, and transform it as needed. Please include some documentation explaining how to run your scripts.

To submit the source code, configuration files, etc, the easiest way is to share a Github or BitBucket repository with us (we will send you the appropriate usernames). Alternatively, we can accept compressed tarballs or zip archives. We cannot accept those over email, though, so we recommend a file sharing service like Google Drive, Dropbox, or similar.

Running your solution

If you use a database engine (like Redshift, PostgreSQL, MySQL, etc), please include instructions for passing in the database connection information (hostname, port, user, password, etc; a single JDBC style url is also fine). If you use another tool for your data warehouse, please include detailed instructions for how to setup that tool. We are experienced developers, but we may not be familiar with the tools or languages you used, so please draft the instructions for running your solution accordingly.

External tools and platforms

Feel free to use whatever external tools, libraries, and platforms you feel are best suited to solve the problem. It is not necessary to write the automation for your solution from the ground up. You can use configuration tools, cloud services, or utility libraries as you see fit.

Assessment and Interview

After we receive your submission we will conduct a review of your solution and execute some tests. We will assess your solution on the following criteria with our review and testing:

- *Correctness*: Does your solution adhere to the specification and does it correctly answer the questions posed above?
- *Scalability*: If we were to scale the game data from its current form to millions of players and tens of millions of games will your solution continue to work without become unusably slow or otherwise broken?
- *Automation*: How easily is your solution executed and changed?
- *Readability*: Can an engineer unfamiliar with your implementation read and understand what you have implemented, and then make changes? This speaks to style, naming conventions, organization, and documentation.

On the day of your on-site interview you will present your solution to 2-3 members of the engineering team. You should prepare to talk about your implementation approach, design trade offs, and approach to testing and validation.

Through the course of the one-on-one interviews we will ask you further questions about how you would extend your solution and how you would fix any issues we find in our own testing and evaluation to improve your solution.

